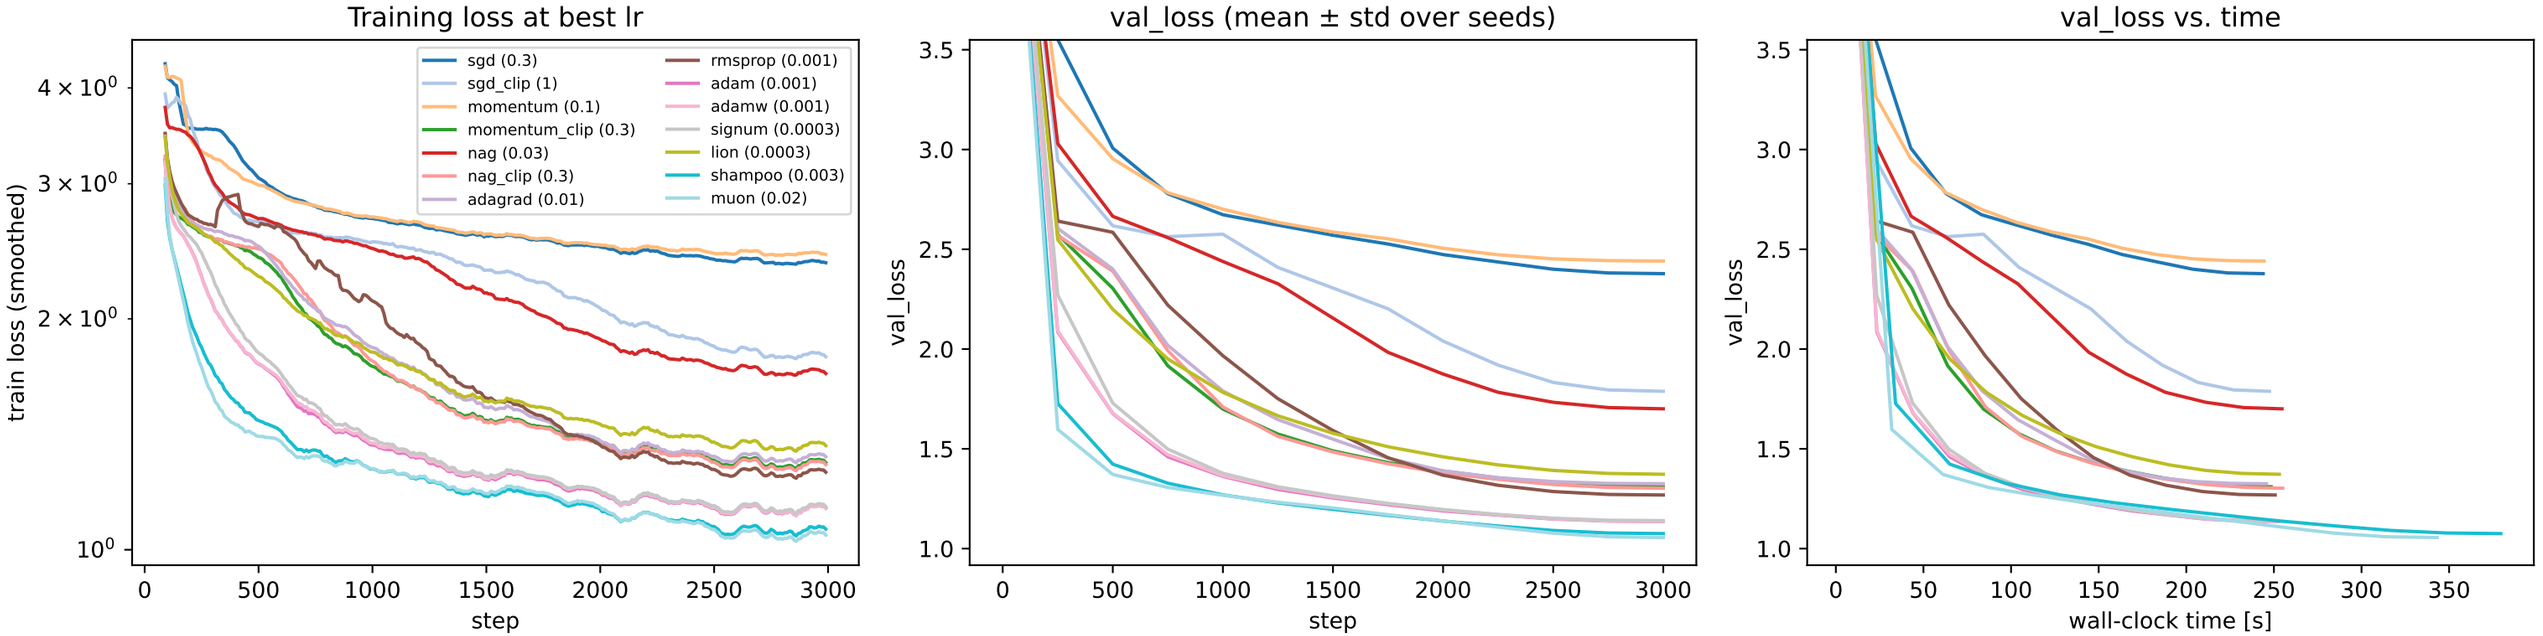
Training loss at best lr
train loss (smoothed)
step
sgd (0.3)
sgd_clip (1)
momentum (0.1)
momentum_clip (0.3)
nag (0.03)
nag_clip (0.3)
adagrad (0.01)
rmsprop (0.001)
adam (0.001)
adamw (0.001)
signum (0.0003)
lion (0.0003)
shampoo (0.003)
muon (0.02)
val_loss (mean ± std over seeds)
val_loss
step
val_loss vs. time
val_loss
wall-clock time [s]