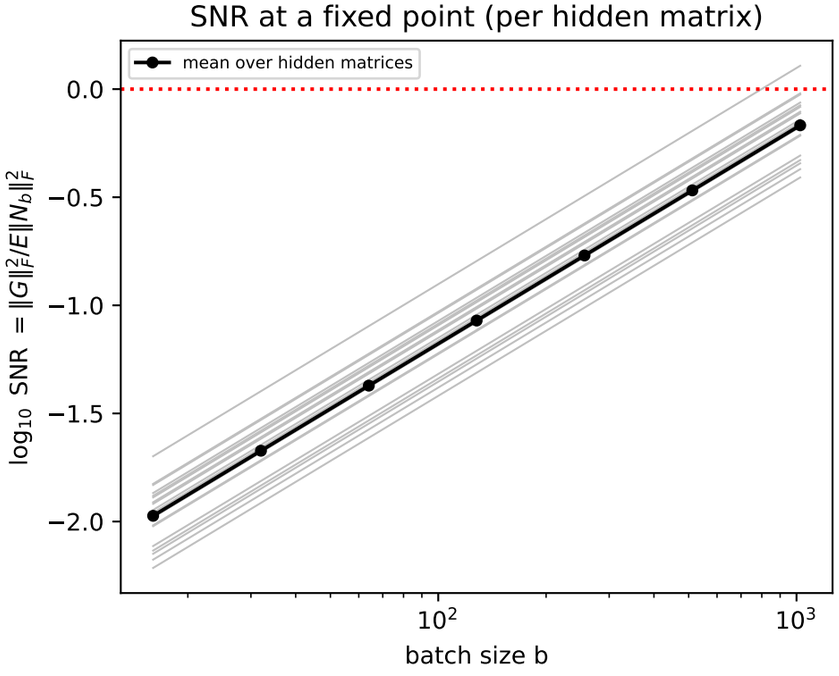
SNR at a fixed point (per hidden matrix)
mean over hidden matrices
$\log_{10}$ SNR $= \|G\|_F^2/E\|N_b\|_F^2$
batch size b

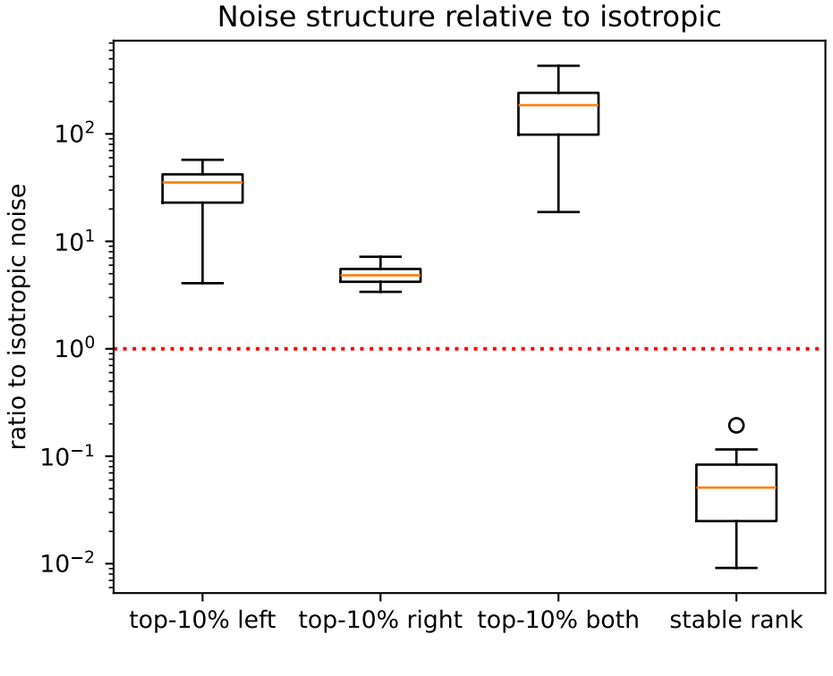
Noise structure relative to isotropic
ratio to isotropic noise
top-10% left
top-10% right
top-10% both
stable rank

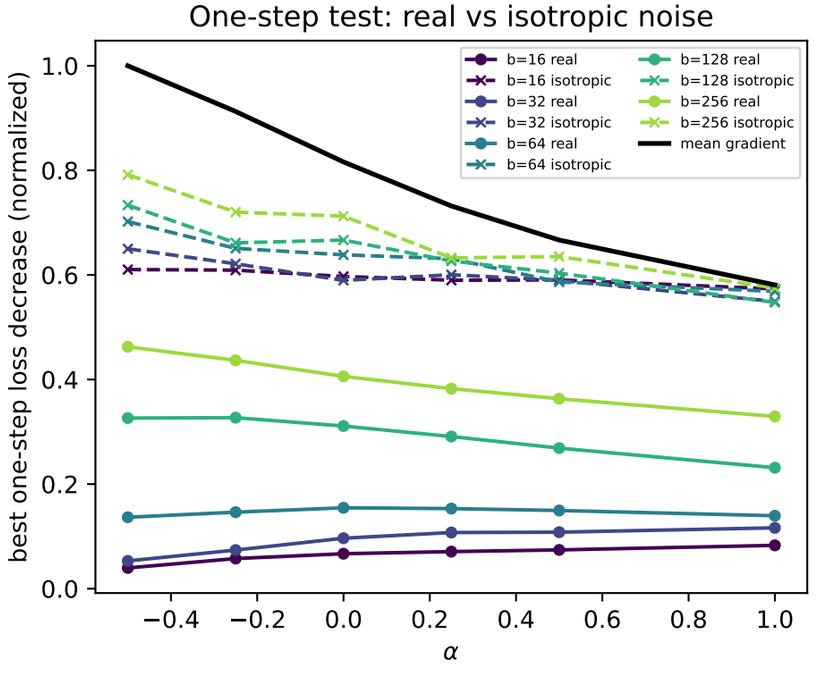
One-step test: real vs isotropic noise
best one-step loss decrease (normalized)
b=16 real
b=16 isotropic
b=32 real
b=32 isotropic
b=64 real
b=64 isotropic
b=128 real
b=128 isotropic
b=256 real
b=256 isotropic
mean gradient
α